

Pointers and Memory in C++



A beginner's guide - by Steven Gabarró

Slides and Examples...

- They can all be found at
 - <https://github.com/bewchabbacc/pointers>
- Note that all examples include “helpers.h” which contains a function printing the address and content of a variable
- The header also contains
 - `#include <iostream>`
 - `using namespace std;`
- Note that the “using” line has been placed in the header to make the examples shorter to read. It is usually a bad idea to include it in the global scope, and should be kept at the file level...

Virtual Memory - A refresher

- Each process has access to its own “virtual memory”, starting at address 0, mapped to physical memory by the OS
- Virtual Memory is typically organized as (this may vary slightly depending on architecture/OS):
 - OS/Architecture-Dependent Region (Off-limits)
 - Text (aka your compiled code)
 - Data (Initialized global and static data)
 - BSS (Block Started by Symbol - Uninitialized static/global data)
 - Heap (or Free memory, growing up)
 - Large empty space (Dynamic libraries are loaded in here)
 - Stack (growing down)
 - OS/Architecture-Dependent Region

The Heap

- The heap, or free store, is an area of memory that is shared throughout a process
- Data (variables / objects) stored in it can be read as long as the address of it is known
- Grows and shrinks by allocating / deallocating data (more on that later)

The Stack

- This is an area of memory that contains your stack frames (or Execution Block, or Activation Record)
- Each function call creates a new stack frame containing
 - The parameters sent to the function
 - The local variables of that function
 - Control data such as a reference to the previous stack frame and the return address
- It grows with each function call, and shrinks on each return

What is a pointer?

- Simply put: It is the address of something
- Pointers are typically the size of registers in your processor (4B in 32bit processors, 8B in 64bit)
- You can have pointers to:
 - Primitive type variables
 - Objects
 - Other pointers
 - Or pointers to pointers
 - Or pointers to pointers to pointers
 - etc...
 - Functions
 - Essentially, anything

Pointer Syntax

- To declare a pointer:
 - `type *var; // or type* var;`
- Ideally, initialize pointers with NULL if unsure about what to store in it
 - `type *var = NULL;`
- You can retrieve the address of any variable with the & symbol
 - `int a = 42;`
 - `int *pi = &a; // pi now holds the address of pi`
- You can access the block pointed to by var by *dereferencing* it
 - `*var = 5;`

Visual Example

- `int var = 5;`
`int *ptr = &var;`

var == 5

ptr



Common Issues

- Do NOT assume pointers get initialized automatically with a 0, so you should ALWAYS give it a value (e.g. NULL (or nullptr in C++11))
- Understand the difference between
 - `type *a, b, c;` AND `type* a, b, c;`
 - In essence, NEVER do `type * a, b, c;`

Proper uses of *

(other than multiplications)

- type *symbol
 - Declaration of a pointer (also used in function parameter declarations)
 - `int *pi = NULL;`
 - `void f(int *pi) { ... }`
- type*
 - Used when declaring (but not implementing) methods
 - `void f(char*);`
- *var
 - Used to dereference a pointer (or reading what the pointer points to)
 - Note that the size read will depend on the type of the pointer!
 - `*pi = 5; // copies 5 into the integer block pointed by pi`
- type (*name) (params...)
 - Function pointers, try to avoid them unless absolutely necessary
 - `int (*operation)(int, int);`

Proper uses of &

- `&var`
 - Get the address of the variable var
- `type& var`
 - Used in function declarations to specify var is sent by reference
 - `void f(int& var) { ... }`
- `var1 & var2`
 - Binary “and” operation between two variables
- `expr1 && expr2`
 - Logical “and” between two expressions

What NOT to do...

- `&&var`
 - This would mean get the address of the address of var
 - In reality it's nonsense!
 - Note: This is not the same as the `&&` used in the C++11 “move” operator
- `*(&var)`
 - That would mean dereference the address of var
 - That's silly, just use var...
- `int *pi = (int*)5;`
 - While technically possible, this would mean that pi is now a pointer to the memory address 5!!! This would result in unpredictable code, and luckily most compilers will warn you about it
- Dereferencing a Bad pointer
- Returning the address of a Stack variable (compilers will complain about it)

```
Executor:pointers bewchy$ make test3
c++ -std=c++03 -Wall -pedantic-errors -c -o test3.o test3.cpp
test3.cpp:8:11: warning: address of stack memory associated with local variable 'j' returned
      [-Wreturn-stack-address]
    return &j; // YIKES!! It hurts me to type this...
               ^
1 warning generated.
c++ -std=c++03 -Wall -pedantic-errors test3.o -o test3.out
```

```
#include "helpers.h"

int *f(int i)
{
    int j = 5;

    j += i;
    return &j; // YIKES!! It hurts me to type this...
}

int main()
{
    int *pi1 = NULL;
    int *pi2 = NULL;

    pi1 = f(6);
    print_info(pi1);
    cout << "pi1 address " << pi1 << " with data " << *pi1 << endl;
    pi2 = f(5);
    print_info(pi2);
    cout << "pi2 holds " << pi2 << " with data " << *pi2 << endl;
    cout << "pi1 holds " << pi1 << " with data " << *pi1 << endl;

    return 0;
}
```

What is a bad pointer?

- Cristobal Colon statue in Barcelona (on the opening slide)
- He should be pointing to the Americas, but instead...



But really: What is a Bad Pointer?

- Pointer pointing to an area that you are not in control of
- Best case scenario, it's an area protected by the OS and you will get a segmentation fault when attempting to dereference the pointer
- Worst case scenario, you just played around with a block of memory that was holding something important needed elsewhere in your code (yikes!)

```

#include "helpers.h"

void f_ptr(int **x)
{
    cout << "x is a pointer sent by value->";
    print_info(x);
}

void f_ptr_ref(int* &x)
{
    cout << "x is a pointer sent by reference->";
    print_info(x);
}

```

```

Executor:pointers bewchy$ g++ -DCRASHME test1.cpp -o test1.out
Executor:pointers bewchy$ ./test1.out
Segmentation fault: 11
Executor:pointers bewchy$ 

```

```

Executor:pointers bewchy$ make test1
c++ -std=c++03 -Wall -pedantic-errors -c -o test1.o tes
c++ -std=c++03 -Wall -pedantic-errors test1.o -o test1.out

##### test1 #####
i->@: 0x7fff5b5b7b28 Sz: 4 Val: 5
pi->@: 0x7fff5b5b7b20 Sz: 8 Val: 0x7fff5b5b7b28
*pi->@: 0x7fff5b5b7b28 Sz: 4 Val: 5
i->@: 0x7fff5b5b7b28 Sz: 4 Val: 8
pi->@: 0x7fff5b5b7b20 Sz: 8 Val: 0x7fff5b5b7b28
*pi->@: 0x7fff5b5b7b28 Sz: 4 Val: 8

```

```

Let's see our other pointers...
@: 0x7fff5b5b7b18 Sz: 8 Val: 0x0
@: 0x7fff5b5b7b10 Sz: 8 Val: 0x0
x is a pointer sent by value->@: 0x7fff5b5b7ab8 Sz: 8 Val: 0x7fff5b5b7b28
x is a pointer sent by reference->@: 0x7fff5b5b7b20 Sz: 8 Val: 0x7fff5b5b7b28

```

```

int main()
{
    int i = 5;
    int *pi = &i;
    int *pi2 = NULL; // Dereferencing this will seg-fault!
    int *pi3; // THIS IS A BAD IDEA!

#ifndef CRASHME
    int *pi4 = (int*)5; // THIS IS EVEN WORSE!!!
#endif

    *pi4 = 10; // Congrats! you just got a segfault! I hope you're proud
#endif

    cout << "i->";
    print_info(i);
    cout << "pi->";
    print_info(pi);
    cout << "*pi->";
    print_info(*pi);
    i = 8;
    cout << "i->";
    print_info(i);
    cout << "pi->";
    print_info(pi);
    cout << "*pi->";
    print_info(*pi);

    cout << "\nLet's see our other pointers..." << endl;
    print_info(pi2);
    print_info(pi3);

    f_ptr(pi);
    f_ptr_ref(pi);
    return 0;
}

```

Beware of pointer arithmetic!

- Note that basic operations may not behave as you expect
- `ptr = ptr2`
 - Copies the data in the pointer, which is an address!
 - Does not copy the data pointed by the pointers
- `ptr++` or `++ptr`
 - Increments the pointer by `sizeof(*ptr)` NOT just 1 (unless it is a `char*`)
 - e.g.
 - `int *ptr = &i; ++ptr;`
 - Now `ptr` is pointing `sizeof(int)` bytes AFTER `i`
- `*(ptr++)` and `(*ptr)++`
 - `*(ptr++)` will dereference the pointer `ptr`, then increment it by `sizeof(*ptr)`
 - `(*ptr)++` will increment the data pointer by the pointer

The `new` and `delete` keywords

- While you can always get the address of something, you typically want the address of something that is NOT in the stack
- In order to put things in the Heap, we use `new`, to remove them we use `delete`
- `int *pi = new int;`
 - This creates a new block that can hold an integer in the Heap, and returns the address to pi
 - While similar to C's `malloc`, they are not interchangeable (especially when allocating objects, as `new` calls constructors but `malloc` doesn't)
- `delete pi;`
 - This marks the Heap block pointed by pi as *free*
 - This is similar to C's `free`, but not interchangeable (this will call object destructors)

EXTREMELY IMPORTANT!

DO NOT *free* SOMETHING ALLOCATED WITH *new*
DO NOT *delete* SOMETHING ALLOCATED WITH *malloc*
DO NOT *free* NOR *delete* THE SAME THING TWICE
DO NOT *free* NOR *delete* A STACK VARIABLE

```
#include "helpers.h"

int *f()
{
    int *ret = new int;

    *ret = 42;
    return ret;
}

int main()
{
    int i = 5;
    int *pi = new int;
    int *pi2 = f();

    cout << "pi was allocated but not initialized" << endl;
    print_info(pi);
    print_info(*pi);
    *pi = i;
    print_info(pi);
    print_info(*pi);
    cout << "Let's checkout pi2: \n";
    print_info(pi2);
    print_info(*pi2);

    delete pi;
    delete pi2;
    return 0;
}
```

```
Executor:pointers bewchy$ make test2
c++ -std=c++03 -Wall -pedantic-errors -c -o test2.o test2.cpp
c++ -std=c++03 -Wall -pedantic-errors test2.o -o test2.out

##### test2 #####
pi was allocated but not initialized
@: 0x7fff5a19eb10 Sz: 8 Val: 0x7fdc8c04a20
@: 0x7fdc8c04a20 Sz: 4 Val: 0
@: 0x7fff5a19eb10 Sz: 8 Val: 0x7fdc8c04a20
@: 0x7fdc8c04a20 Sz: 4 Val: 5
Let's checkout pi2:
@: 0x7fff5a19eb08 Sz: 8 Val: 0x7fdc8c04a30
@: 0x7fdc8c04a30 Sz: 4 Val: 42
```

BEWARE!

- If you do not free a pointer and lose track of where it is, you just caused a *memory leak*
- This means you now have a reserved block of memory in the heap you have lost track of, taking up space for no reason
- e.g.
 - `int i = 42; // makes a local block`
 - `int *pi = new int; // allocates a heap block`
 - `pi = &i; // overwrites the heap pointer with the address of the local i`

Basic arrays

- Arrays are a contiguous block of memory holding data of a certain type
- The “name” of the array is in fact a pointer to the first byte of that block
- When you do `arr[i]` it is the same as doing
 - `*(&arr + (i * sizeof(*arr)))`
 - or `(arr+(i*sizeof(arr[0])))`
- This means that it is technically legal to do `arr[-1]` but **DO NOT DO IT!**

delete[] and new[]

- In order to allocate a “raw” array of type T, you can do
 - `T arr[n]; // this puts the array in the stack`
 - `T arr[] = {data1, data2, data3, ...};`
 - `T *arr2 = new T[n]; // this puts the array in the heap`
- Only put arrays in the heap if they may be needed outside the scope of your function
 - Better yet, use containers (see later on)
- To delete an array from the Heap, use `delete []`
 - `delete [] arr2;`
- Do NOT delete [] stack arrays
- Do NOT forget the [] when deleting arrays!

```
#include "helpers.h"

int main()
{
    int arr[] = {1, 2, 3, 4, 5};
    int arr2[5];
    int *arr3 = new int[5];

    for (int i = 0; i < 5; ++i)
    {
        arr2[i] = 2*i;
        arr3[i] = 3*i;
    }
    cout << "Arr info: ";
    print_info(arr);
    cout << "\nArr2 info: ";
    print_info(arr2);
    cout << "\nArr3 info: ";
    print_info(arr3);
    cout << "\nArr[1]: ";
    print_info(arr[1]);
    cout << "\nArr2[2]: ";
    print_info(arr2[2]);
    cout << "\nArr3[3]: ";
    print_info(arr3[3]);
    cout << endl;

    delete [] arr3;
    return 0;
}
```

```
Executor:pointers bewchy$ make test4
c++ -std=c++03 -Wall -pedantic-errors -c -o test4.o test4.cpp
c++ -std=c++03 -Wall -pedantic-errors test4.o -o test4.out

##### test4 #####
Arr info: @: 0xffff55831ab0 Sz: 20 Val: 0xffff55831ab0
Arr2 info: @: 0xffff55831a90 Sz: 20 Val: 0xffff55831a90
Arr3 info: @: 0xffff55831a70 Sz: 8 Val: 0x7fb04ac04a20
Arr[1]: @: 0xffff55831ab4 Sz: 4 Val: 2
Arr2[2]: @: 0xffff55831a98 Sz: 4 Val: 4
Arr3[3]: @: 0x7fb04ac04a2c Sz: 4 Val: 9
```

Allocating Objects: constructors and destructors

- When you create an object from a class, its constructor gets called
- When it is destroyed, the destructor is called
- If your object is in the Stack, constructor is called when you create it, the destructor is called when you “fall out of scope”
- If you want your object to be in the Heap, constructor/destructor will be called when you do new/delete
- To access data members and methods of a Stack object do
 - `obj.method()`
- To access data members and methods through a pointer to an object you could do
 - `(*obj).method() // UGLY!`
 - `obj->method() // NICER!`

```

#include "helpers.h"
#include <stdexcept>

class IntArrWrapper
{
private:
    int cap;
    int *data;
public:
    IntArrWrapper() {cap = 5; data = new int[5];}
    IntArrWrapper(int c) {
        if (c < 5) // we're forcing a minimum of 5 ints (why not)
            c = 5;
        cap = c;
        data = new int[c];
    }
    ~IntArrWrapper() { delete [] data; cout << "Deleted the data!" << endl; }
    int& operator[](int n) throw (out_of_range) {
        if (n < 0 || n >= cap)
            throw out_of_range("Illegal index!\n");
        return data[n];
    }
    int length() const {
        return cap;
    }

    friend ostream& operator<< (ostream&, const IntArrWrapper&);
};

ostream& operator<< (ostream& o, const IntArrWrapper& a)
{
    o << "\nCap var at: " << &(a.cap);
    o << "\ndata at " << &(a.data) << " its pointer is " << a.data << endl;
    return o;
}

```

```

int main() {
    IntArrWrapper arr;
    IntArrWrapper *arr2 = new IntArrWrapper(10);

    cout << "ARR:\n";
    print_info(arr);
    cout << "ARR2:\n";
    print_info(arr2);
    cout << "*ARR2:\n";
    print_info(*arr2);
    arr[1] = 42;
    print_info(arr[1]);

    delete arr2;
    return 0;
}

##### test5 #####
ARR:
@: 0x7fff5b927b28 Sz: 16 Val:
Cap var at: 0x7fff5b927b28
data at 0x7fff5b927b30 its pointer is 0x7fdd73500000

ARR2:
@: 0x7fff5b927b20 Sz: 8 Val: 0x7fdd73500020
*ARR2:
@: 0x7fdd73500020 Sz: 16 Val:
Cap var at: 0x7fdd73500020
data at 0x7fdd73500028 its pointer is 0x7fdd73500030

@: 0x7fdd73500004 Sz: 4 Val: 42
Deleted the data!
Deleted the data!

```

Inheritance and Object Allocation

- If you are using inheritance be aware of constructor and destructor call order
- The topmost constructor in the class hierarchy is called first, working its way down to the object you constructed
- Destructors are called in reverse order (from lowest in the hierarchy, working up through the ancestors)

```

class Base
{
private:
    static unsigned int baseCount;
protected:
    vector<string> data;
public:
    virtual ~Base()
    {
        --baseCount;
        cout << "Base Destructor! EXTERMINATE!!" << endl;
    }
    Base()
    {
        ++baseCount;
        cout << "Base Constructor!" << endl;
    }
    Base(unsigned int n)
    {
        ++baseCount;
        cout << "Base non-default constructor!" << endl;
        data.reserve(n);
    }
    void add(string s)
    {
        data.push_back(s);
    }
    static unsigned int classCount() {
        return baseCount;
    }
    virtual void print_data()
    {
        for (vector<string>::iterator it = data.begin(); it != data.end(); ++it)
        {
            cout << *it << " ";
        }
        cout << "\nThis Base object is at location: " << this << endl;
    }
};

```

```

class Child: public Base
{
private:
    static unsigned int childCount;
    unsigned int count;
public:
    ~Child() {
        --childCount;
        cout << "Child destructor!" << endl;
    }
    Child() {
        ++childCount;
        count = 0;
        cout << "Child default constructor!" << endl;
    }
    Child(unsigned int n): Base(n)
    {
        ++childCount;
        count = n;
        cout << "Child non-default constructor!" << endl;
    }
    void print_data()
    {
        cout << "There are " << count << " items: ";
        Base::print_data();
        cout << "This Child object is at location: " << this << endl;
    }
    static unsigned int classCount() {
        return childCount;
    }
};

unsigned int Base::baseCount = 0;
unsigned int Child::childCount = 0;

```

```

int main(int argc, char **argv)
{
    if (argc < 2)
        cout << "Use: ./test3.out string1 string2 string3...";
    Child c(argc-1);
    Base *b = &c;
    Child *c2 = new Child(argc - 1);
    for (int i = 1; i < argc; i++)
    {
        c.add(string(argv[i]));
        c2->add(string(argv[i]));
    }
    cout << endl << "Showing contents through Child class" << endl;
    c.print_data();
    cout << endl << "Showing contents through Child pointer class" << endl;
    c2->print_data();
    cout << endl << "Showing contents through base class\n" << endl;
    b->print_data();

    cout << "\nThere are " << Base::classCount() << " base objects made and ";
    cout << Child::classCount() << " children ones." << endl;

    delete c2;
    return 0;
}

```

```

Executor:pointers bewchy$ make test6
c++ -std=c++03 -Wall -pedantic-errors -c -o test6.o test6.cpp
c++ -std=c++03 -Wall -pedantic-errors test6.o -o test6.out

```

```

##### test6 #####
Base non-default constructor!
Child non-default constructor!
Base non-default constructor!
Child non-default constructor!

```

```

Showing contents through Child class
There are 9 items: This is a call to all my past resignations
This Base object is at location: 0x7fff5d4b19a8
This Child object is at location: 0x7fff5d4b19a8

```

```

Showing contents through Child pointer class
There are 9 items: This is a call to all my past resignations
This Base object is at location: 0x7fada0404b20
This Child object is at location: 0x7fada0404b20

```

```

Showing contents through base class

```

```

There are 9 items: This is a call to all my past resignations
This Base object is at location: 0x7fff5d4b19a8
This Child object is at location: 0x7fff5d4b19a8

```

```

There are 2 base objects made and 2 children ones.
Child destructor!
Base Destructor! EXTERMINATE!!
Child destructor!
Base Destructor! EXTERMINATE!!

```

Naked pointers vs containers

- Until now we have talked about “naked pointers” aka “old school pointers”
- IF you chose to use naked pointers and allocate space yourself, remember that it is YOUR job to clean it up (collect your own garbage!)
- If you are trying to store data in memory, avoid using raw pointers / arrays. Use containers instead!
 - They handle memory themselves
 - Will grow/shrink as needed
 - Will get deleted for you (unless you allocated a pointer to a container, in which case you still need *delete*)
 - So PLEASE use things like *vector*, *deque*, *list*, *stack*, *queue*, *priority_queue*, *set*, *multiset*, *map*, or *multimap*

More Pointers Danger

- NEVER *delete* a stack variable / object / array
- NEVER *delete* something that is not a pointer
- NEVER assume pointers are initialized for you
- If you own a pointer, *delete* it when done with it! This includes pointers you did not allocate but were created for you by other functions
- Watch out for pointer arithmetic errors!
- If you pass a pointer to a function, by default you are sending a copy of that pointer
- If you have the choice between passing data by reference or passing a pointer to your data, send by reference!
- If you make a copy of a pointer, don't delete it twice! *delete* gets rid of the pointee, not the pointer!
- If you need data in the Heap, consider using a container or make a class that will take care of the data with a proper constructor and destructor

Quiz! Find the bugs!

- arr is a copy of the pointer ptr, so changes to arr don't affect ptr
- Adding one to a pointer will place your pointer 1B into the first number, rather than going to the next number in the array
- There's a memory leak caused by the allocation of ptr followed by the overwriting of the pointer with the data array
- The delete is wrong!
 - It has nothing to do with delete []
 - Since you overwrote ptr with data, now ptr refers to a stack array so DON'T DELETE IT!
- We need a return statement in main!

```
#include "helpers.h"

// Supposedly grabs the item pointer by arr, and
// moves arr up by one item
int each(int *arr)
{
    int ret = *arr;

    arr += 1;
    return ret;
}

int main()
{
    int data[] = {32, 54, 33, 76, 4, 26, 76};
    int *ptr = new int[7];

    ptr = data;
    for (int i = 0; i < 7; ++i)
        cout << each(ptr) << " ";
    delete ptr;
}
```

A peek into the “future”

- To avoid many of the problems you can get from using raw pointers, C++11 introduced “smart pointers”
 - `unique_ptr<type>`
 - Prevents the pointer from being copied
 - `shared_ptr<type>`
 - Only gets deleted with ALL copies of the pointer are no longer needed/used
 - `weak_ptr<type>`
 - Created as a copy of a `shared_ptr` that won’t affect the other `shared_ptr`, but will get erased when the `shared_ptr` it came from gets deleted
- Note that C++03 has *auto_ptr*
 - This was removed in C++11 and in order to future-proof your code, you should avoid using it!

Questions?